

Thai Sign Language — Raw Video Collection


Harvest report · generated 2026-09-13 04:18 UTC+0700

5,331 video clips	25h 34m 50s total footage	11.8 GB on disk	5,035 distinct labels
-----------------------------	-------------------------------------	---------------------------	---------------------------------

Every clip in this collection is raw video — **.mp4 as downloaded**, no pose extraction, no feature files. Each one carries a text label (a Thai gloss, or the publisher's own title) so the corpus is usable both for self-supervised pre-training and for downstream supervised work.

1 What was collected

Level	Clips	Footage	Size	Folder
word	5,229	16h 49m 21s	8.3 GB	raw_data/word_level/
sentence	57	2h 36m 21s	1.2 GB	raw_data/sentence_level/
continuous	45	6h 9m 7s	2.3 GB	raw_data/continuous/
total	5,331	25h 34m 50s	11.8 GB	



2 Yield by source

Source	Clips	Footage	Size
ttrs_dictionary	5,043	6h 10m 34s	3.9 GB
youtube	288	19h 24m 15s	7.9 GB

3 Sources attempted, and their status

Provenance is part of the deliverable. Two of the requested sources were deliberately left alone, and two were unreachable — the reasons are recorded here rather than buried in a log.

Source	Status	Detail
Thai Telecommunication Relay Service (TTRS), NECTEC/NSTDA https://dic.ttrs.or.th	HARVESTED	Primary word-level corpus. Public JSON API used by the site's own front-end. no robots.txt served; no crawl restriction declared
Thai PBS -- 'Big Sign' full-frame sign language https://www.youtube.com/channel/UCFpFQ_j7AbBjD0Dpc_ZJQqQ	HARVESTED	Full-screen Thai Sign Language programming; signer fills the frame. accessed through yt-dlp with request throttling
Mixed Thai publishers, schools, deaf-community channels YouTube search (9 curated Thai-language queries)	HARVESTED	Titles are filtered for an explicit sign-language marker so that every clip carries a usable text label. accessed through yt-dlp with request throttling
National Association of the Deaf in Thailand (สมาคมคนหูหนวกแห่งประเทศไทย) https://www.th-sl.com	NOT HARVESTED -- express reservation of rights	The site's robots.txt both disallows this agent by name and carries an express machine-readable reservation against AI training. Excluded. A research licence should be requested directly from NADT. Content-Signal: ai-train=no, use=reference; User-agent: ClaudeBot -> Disallow: /
Royal Society of Thailand (สำนักงานราชบัณฑิตยสภา) https://royalsociety.go.th	NOT HARVESTED -- no video resource located	The Royal Society publishes sign-language terminology in print/PDF form. No sign-language video archive was found on its web properties. -
Parliament of Thailand / TPchannel https://www.parliament.go.th , TPchannel on YouTube	PARTIAL -- interpreter-inset material only	Full session broadcasts are multi-hour with a small interpreter inset, so they need cropping before they are useful. Only briefing-style clips matching the sign-language filter were taken. accessed through yt-dlp with request throttling
TOMTUN / ThaiSignVis thaisignvis.org (DNS does not resolve); tomtun.com (HTTP 520)	NOT HARVESTED -- host unreachable	Neither host resolved to a working sign-language video service at collection time. No public mirror or dataset release was located. -
Kasetsart University -- ThSL digital dictionary https://pirun.ku.ac.th/~fhumalt/THSL/	NOT HARVESTED -- obsolete media format	The dictionary serves its signs through Adobe Shockwave/Flash, which no longer runs and exposes no extractable video files. -

4 Label coverage

Label type	Clips	What it is
gloss	5,043	Curated Thai gloss from the TTRS dictionary, with definition, part of speech and category alongside.
title	288	Publisher's own video title, plus description and any subtitle track that came with it.

Top vocabulary categories (TTRS)

Category (Thai)	Clips
คำในบทเรียน	402
COVID - 19	273
คำกริยาวลี	188



Category (Thai)	Clips
การกระทำ	168
สัตว์	144
ผัก	142
คำนาม	136
สถานที่	107
คำวิเศษณ์	106
อาหาร	105
อาชีพ	104
ห้องครัว	96
คำศัพท์เกี่ยวกับศาล	94
ลักษณะผู้คน	87

136 distinct categories in total.

5 Technical profile

Relevant to SignDINO-style pre-training: these are the raw decode characteristics of the corpus as delivered.

Statistic	Value
clips with a readable duration	5,331
median clip length	4.2 s
p10 / p90 clip length	2.9 s / 7.1 s
shortest / longest	1.8 s / 54m 2s
mean bitrate	1097 kbps



Frame rate and codec

Frame rate	Clips	Video codec	Clips
25 fps	4,995	h264	5,299
30 fps	195	vp9	23
50 fps	101	av1	9
60 fps	19		
24 fps	14		

Frame rate	Clips	Video codec	Clips
15 fps	4		

6 Collection reliability

Outcome	Items	Share of catalogue
done	5,331	86.8%
pending	806	13.1%
failed	2	0.0%
catalogued	6,139	100.0%

Most common failure modes

Source	Error	Count
youtube	skipped: YouTube bot-block breaker is open	789
youtube	ERROR: [youtube] 3flvMCMXQy0: Please sign in. Use --cookies-from-browser or --cookies for the authentication. See https://github.com/yt-dlp/yt-dlp/wi	1
youtube	ERROR: [youtube] 4PowomwflD8: Sign in to confirm you're not a bot. Use --cookies-from-browser or --cookies for the authentication. See https://github.com/yt-dlp/yt-dlp/wi	1
youtube	ERROR: [youtube] B2cUKdtSMbl: Sign in to confirm you're not a bot. Use --cookies-from-browser or --cookies for the authentication. See https://github.com/yt-dlp/yt-dlp/wi	1
youtube	ERROR: [youtube] CEsjVl1ihc: Please sign in. Use --cookies-from-browser or --cookies for the authentication. See https://github.com/yt-dlp/yt-dlp/wi	1
youtube	ERROR: [youtube] lh2iN9esDnw: Sign in to confirm you're not a bot. Use --cookies-from-browser or --cookies for the authentication. See https://github.com/yt-dlp/yt-dlp/wi	1
youtube	ERROR: [youtube] MEYYuA2D5wA: Please sign in. Use --cookies-from-browser or --cookies for the authentication. See https://github.com/yt-dlp/yt-dlp/wi	1
youtube	ERROR: [youtube] NerQlXpWsh0: Sign in to confirm you're not a bot. Use --cookies-from-browser or --cookies for the authentication. See https://github.com/yt-dlp/yt-dlp/wi	1
youtube	ERROR: [youtube] S_DWYTcVwg4: Please sign in. Use --cookies-from-browser or --cookies for the authentication. See https://github.com/yt-dlp/yt-dlp/wi	1
youtube	ERROR: [youtube] Ue4w9PESy_Y: Sign in to confirm you're not a bot. Use --cookies-from-browser or --cookies for the authentication. See https://github.com/yt-dlp/yt-dlp/wi	1
youtube	ERROR: [youtube] UwUzSxdK61g: Sign in to confirm you're not a bot. Use --cookies-from-browser or --cookies for the authentication. See https://github.com/yt-dlp/yt-dlp/wi	1
youtube	ERROR: [youtube] cqpQfl-gA6l: Sign in to confirm you're not a bot. Use --cookies-from-browser or --cookies for the authentication. See https://github.com/yt-dlp/yt-dlp/wi	1

How the collector stayed polite and stable

- **Per-host rate limiting.** Every request passes a token bucket keyed on hostname — one request per configured interval plus random jitter, shared across all worker threads, so raising concurrency never raises request rate.
- **Exponential backoff with full jitter.** Retryable failures (408, 425, 429, 5xx, timeouts, connection resets) back off over a window that doubles per attempt, sampled uniformly to avoid retry convoys. A Retry-After header always wins over the computed delay.

- **Adaptive braking.** A 429 or 503 adds a decaying per-host cooldown on top of the static spacing, so a strained server gets more room automatically.
- **Atomic writes.** Downloads stream to .part and are renamed into place only after length, sha256 and an ffmpeg decode check all pass. A truncated file can never be mistaken for a good one.
- **Idempotent resume.** All state lives in a SQLite catalogue. Interrupting the run costs only the in-flight downloads; re-running continues from there and re-downloads nothing.
- **Content-level deduplication.** Clips are hashed; the same bytes reached through two different listings are recorded once.

7 Files you now have

Path	What it holds
raw_data/word_level/	Isolated lexical signs — one gloss per clip.
raw_data/sentence_level/	Single utterances and short phrases.
raw_data/continuous/	Running signing: interpreted news, songs, briefings.
raw_data/_quarantine/	Files that downloaded but failed decode validation.
metadata/metadata_master.csv	Every clip, one row each, all label fields.
metadata/metadata_word.csv ...	The same split by linguistic level.
metadata/failures.csv	Anything not retrieved, with the reason.
metadata/sources_manifest.csv	Provenance and access status per source.
state/harvest.db	SQLite catalogue — the resume point.
*.json (beside each clip)	Lossless per-clip record from upstream.
*.vtt (beside YouTube clips)	Subtitle tracks, where the publisher had them.

8 Before you train on this

- **Licensing is not settled by collection.** These clips were gathered from public endpoints for research. They are not public domain. Clearing redistribution or model release — particularly for the TTRS dictionary and any broadcaster material — is a separate step.
- **The NADT database was deliberately skipped.** th-sl.com carries an express machine-readable reservation against AI training and disallows this agent by name. It is the single richest Thai word-level source, so a direct research licence request to the association is the highest-value next move.
- **Continuous clips need cropping.** Interpreted broadcast material often puts the signer in a corner inset. Big Sign material does not — the signer fills the frame — which makes it the better continuous subset to start from.
- **Titles are weak labels.** YouTube-sourced rows carry the publisher's title, which names the topic rather than glossing each sign. Treat them as coarse supervision; the TTRS rows are the ones with true per-clip glosses.